

Supervised Learning - Shrinkage Methods

Dr. Sameera Viswakula

June, 2021

Shrinkage methods

As an alternative to selecting a subset of predictors using a least squares fit, we can fit a model containing all p predictors using a technique that ***constrains*** or ***regularizes*** the coefficient estimates. This approach:

- ▶ improves the fit
- ▶ reduce the variance of the estimates

The two best-known techniques for shrinking the regression coefficients towards zero are ***ridge regression*** and the ***lasso***.

The Ridge Regression

Recall:

In least squares fitting procedure, we estimate $\beta_0, \beta_1, \dots, \beta_p$ by minimizing RSS, where

$$RSS = \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2$$

Ridge regression is very similar to least squares, except that the coefficients are estimated by minimizing a slightly different quantity. Here we minimize

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2 = RSS + \lambda \sum_{j=1}^p \beta_j^2$$

where $\lambda \geq 0$ is a **tuning** parameter.

The Ridge Regression ..

Here $\lambda \sum_j \beta_j^2$ is called a ***shrinkage penalty***. This penalty is small when $\beta_1, \beta_2, \dots, \beta_p$ are close to zero.

- ▶ when $\lambda = 0$, ridge regression produces least squares estimates.
- ▶ as $\lambda \rightarrow 0$ the impact of the shrinkage penalty grows, and the ridge regression coefficient estimates will approach zero.

Recall that least squares method generates only one set of coefficient estimates. Here in ridge regression, it produces a different set of coefficient estimates, β^R , for each value of λ . Selecting a good value for λ is critical.

The shrinkage penalty is applied to β_1, \dots, β_p , but not to the intercept β_0 .

The Lasso

- ▶ In the ridge regression, all p predictors are in the final model.
- ▶ The penalty, $\lambda \sum_j \beta_j^2$ shrinks the coefficients towards zero.
- ▶ But not exactly to zero.
- ▶ This leads to model interpretations difficulties when p is quite large.
- ▶ Lasso overcomes this disadvantage.

The lasso coefficients, $\hat{\beta}_\lambda^L$, minimize the quantity

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p |\beta_j| = RSS + \lambda \sum_{j=1}^p |\beta_j|$$

Lasso uses l_1 penalty. l_1 norm is defined as $||\beta||_1 = \sum |\beta_j|$.

- ▶ l_1 penalty has the effect of forcing some of the coefficient estimates to be exactly equal to zero when the tuning parameter λ is sufficiently large.
- ▶ Hence, the lasso performs variable selection.

Example: Ridge Regression

Here we use the **Hitters** data in the **ISLR** package. We want to predict a baseball player's salary on the basis of various statistics associated with performance in the previous year.

```
library (ISLR)
```

Warning: package 'ISLR' was built under R version 4.0.5

```
data(Hitters)
```

- ▶ What are the variable names?
- ▶ How many records, variables are there?
- ▶ How many missing values are there in the *Salary* variable?

```
# removing missing values
```

```
myHitters = na.omit(Hitters )
```

Example: Ridge Regression..

Here we use the *glmnet* package and its *glmnet()* function to fit ridge regression models.

```
x = model.matrix(Salary~.,myHitters )[, -1]  
y = myHitters$Salary
```

The *glmnet()* function has an ***alpha*** argument that determines what type of model is fit. If ***alpha=0*** then a ridge regression model is fit, and if ***alpha=1*** then a lasso model is fit.

```
library(glmnet)
```

Warning: package 'glmnet' was built under R version 4.0.5

Loading required package: Matrix

Loaded glmnet 4.1-1

Example: Ridge Regression..

```
ridge.mod = glmnet(x,y,alpha=0)
ridge.mod
```

Call: glmnet(x = x, y = y, alpha = 0)

	Df	%Dev	Lambda
1	19	0.00	255300
2	19	1.16	232600
3	19	1.27	211900
4	19	1.39	193100
5	19	1.53	176000
6	19	1.67	160300
7	19	1.83	146100
8	19	2.00	133100
9	19	2.19	121300
10	19	2.39	110500
11	19	2.61	100700
12	19	2.85	91740

Example: Ridge Regression..glmnet()

Defaults of ***glmnet()*** function:

- ▶ performs ridge regression for an automatically selected range of λ values.
- ▶ standardizes the variables so that they are on the same scale.
- ▶ stores regression coefficients for each value of λ in a matrix.
- ▶ estimated coefficients can be accessed by *coef()*.

Example: Ridge Regression..glmnet()

For this example we will be using a grid of λ values ranging from $\lambda = 10^{10}$ to $\lambda = 10^{-2}$.

```
grid = 10^seq(10,-2, length=100)
ridge.mod =glmnet(x, y, alpha=0, lambda=grid)
```

Due to the l_2 norm, the coefficient estimates are much smaller, for large value of λ , as compared to a small value of λ is used.

Example: Ridge Regression..

```
ridge.mod$lambda [50]
```

```
[1] 11497.57
```

```
coef(ridge.mod)[,50]
```

(Intercept)	AtBat	Hits	HmRun	RBI	Walks	Years	CAtBat	CHmRun	CRuns	CRBI	CWalks	DivisionW	PutOuts	Assists	Errors
407.356050200	0.036957182	0.138180344	0.524629976	0.239841459	0.289618741	1.107702929	0.003131815	0.087545670	0.023379882	0.024138320	0.025015421	-6.215440973	0.016482577	0.002612988	-0.020502690

```
sqrt(sum(coef(ridge.mod)[-1,50]^2) )
```

```
[1] 6.360612
```

Example: Ridge Regression..

```
ridge.mod$lambda[60]
```

```
[1] 705.4802
```

```
coef(ridge.mod)[,60]
```

(Intercept)	AtBat	Hits	HmRun	
54.32519950	0.11211115	0.65622409	1.17980910	0.937
Walks	Years	CAtBat	CHits	0
1.31987948	2.59640425	0.01083413	0.04674557	0.337
CRBI	CWalks	LeagueN	DivisionW	Pu
0.09780402	0.07189612	13.68370191	-54.65877750	0.118
Errors	NewLeagueN			
-0.70358655	8.61181213			

```
sqrt(sum(coef(ridge.mod)[-1,60]^2) )
```

```
[1] 57.11001
```

Example: Ridge Regression..predict()

We can use the predict() function to obtain the ridge regression coefficients for a new value of λ , say 50.

```
predict(ridge.mod, s=50, type="coefficients")[1:20 ,]
```

(Intercept)	AtBat	Hits	HmRun	
4.876610e+01	-3.580999e-01	1.969359e+00	-1.278248e+00	1
RBI	Walks	Years	CAtBat	
8.038292e-01	2.716186e+00	-6.218319e+00	5.447837e-03	1
CHmRun	CRuns	CRBI	CWalks	
6.244860e-01	2.214985e-01	2.186914e-01	-1.500245e-01	4
DivisionW	PutOuts	Assists	Errors	
-1.182011e+02	2.502322e-01	1.215665e-01	-3.278600e+00	-9

Example: Splitting the dataset

Let's first set a random seed so that the results obtained will be reproducible.

```
set.seed (1)
train = sample (1: nrow(x), nrow(x)/2)
test = (- train )
y.test = y[test]
```

Next we fit a ridge regression model on the training set, and evaluate its MSE on the test set, using $\lambda = 4$.

```
ridge.mod = glmnet (x[train ,], y[train], alpha=0, lambda=g
thresh = 1e-12)
ridge.pred = predict (ridge.mod, s=4, newx=x[test,])
mean((ridge.pred-y.test)^2)
```

```
[1] 142199.2
```

Example: Ridge Regression..

The test MSE is 101036.8. Now let's fit a model with just an intercept. Note that here the predicted value for each test observation is the mean of the training dataset.

```
mean((mean(y[train ])-y.test)^2)
```

```
[1] 224669.9
```

We could also get the same result by fitting a ridge regression model with a very large value of λ .

```
ridge.pred = predict(ridge.mod, s=1e10, newx=x[test,])  
mean((ridge.pred-y.test)^2)
```

```
[1] 224669.8
```

Example: Ridge Regression..

##Which model is better?

Intercept only model or the ridge regression model with $\lambda = 4$.

Ex: Check whether there is any benefit to performing ridge regression with $\lambda = 4$ instead of just performing least squares regression. Recall that least squares is simply ridge regression with $\lambda = 0$.

Example: Ridge Regression..

Answer

```
ridge.pred = predict (ridge.mod, s=0, newx=x[test ,])  
mean((ridge.pred-y.test)^2)
```

```
[1] 167789.8
```

Example: Ridge Regression..

Comparing estimates

```
lm(y~x, subset=train)
```

Call:

```
lm(formula = y ~ x, subset = train)
```

Coefficients:

(Intercept)	xAtBat	xHits	xHmRun	x
274.0145	-0.3521	-1.6377	5.8145	1
xWalks	xYears	xCatBat	xCHits	xCF
3.7287	-16.3773	-0.6412	3.1632	3
xCRBI	xCWalks	xLeagueN	xDivisionW	xPut
-0.6005	0.3379	119.1486	-144.0831	0
xErrors	xNewLeagueN			
-4.7128	-71.0951			

Example: Ridge Regression..

Comparing estimates

```
predict(ridge.mod, s=0, type="coefficients") [1:20 ,]
```

(Intercept)	AtBat	Hits	HmRun	
274.2089049	-0.3699455	-1.5370022	5.9129307	1.48
Walks	Years	CAtBat	CHits	0
3.7577989	-16.5600387	-0.6313336	3.1115575	3.32
CRBI	CWalks	LeagueN	DivisionW	Pr
-0.5694414	0.3300136	118.4000592	-144.2867510	0.19
Errors	NewLeagueN			
-4.6833775	-70.1616132			

Example: Ridge Regression...Selecting λ


How do we select λ ?

Instead of arbitrarily choosing $\lambda = 4$, it would be better to use cross-validation to choose the tuning parameter. This can be easily done using `cv.glmnet()` function. By default, `cv.glmnet()` function performs ten-fold cross-validation.

```
set.seed (1)
cv.out = cv.glmnet (x[train,],y[train],alpha=0)
```

Example: Ridge Regression.. Cross-validation

```
plot(cv.out)
```



Example: Ridge Regression.. Cross-validation

```
bestlam = cv.out$lambda.min  
bestlam
```

```
[1] 326.0828
```

Therefore, the value of λ that results in the smallest cross-validation error is 211.7416.

What is the test MSE associated with this value of λ ?

```
ridge.pred = predict (ridge.mod, s=bestlam, newx=x[test ,])  
mean((ridge.pred-y.test)^2)
```

```
[1] 139856.6
```

Example: Ridge Regression..

This means we have improved the test MSE.

Finally, we refit our ridge regression model on the full data set, using the value of λ chosen by cross-validation, and examine the coefficient estimates.

```
out = glmnet (x,y,alpha =0)
predict(out, type="coefficients",s=bestlam )[1:20 ,]
```

(Intercept)	AtBat	Hits	HmRun	
15.44383135	0.07715547	0.85911581	0.60103107	1.063
Walks	Years	CAtBat	CHits	0
1.62444616	1.35254780	0.01134999	0.05746654	0.406
CRBI	CWalks	LeagueN	DivisionW	Pu
0.12116504	0.05299202	22.09143189	-79.04032637	0.166
Errors	NewLeagueN			
-1.36092945	9.12487767			

Example: Ridge Regression..Done

As expected, none of the coefficients are zero.

Ridge regression does not perform **variable selection**!

Example: The Lasso

We saw that ridge regression with a wise choice of λ can outperform least squares as well as the null model on the *Hitters* data set.


Can the lasso yield either a more accurate or a more interpretable model than ridge regression?

We use the ***glmnet()*** function with ***alpha=1***.

```
lasso.mod = glmnet(x[train,],y[train],alpha=1, lambda=grid)
plot(lasso.mod)
```

Example: Lasso ..

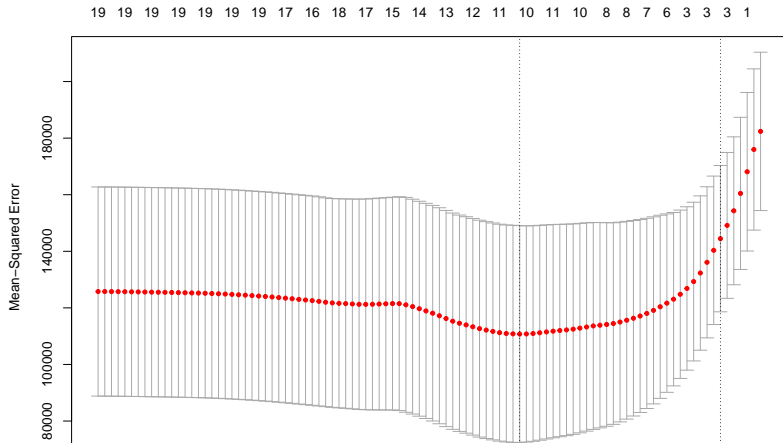
Warning in regularize.values(x, y, ties, missing(ties), na
collapsing to unique 'x' values



Example: Lasso ..

Let's perform cross-validation and compute the associated test error.

```
set.seed (1)
cv.out = cv.glmnet(x[train,], y[train], alpha=1)
plot(cv.out)
```



Example: Lasso ..

```
bestlam = cv.out$lambda.min  
bestlam
```

```
[1] 9.286955
```

```
lasso.pred = predict(lasso.mod, s=bestlam, newx=x[test ,])  
mean((lasso.pred-y.test)^2)
```

```
[1] 143673.6
```

Example: Lasso ..

Estimates

```
out = glmnet (x, y, alpha=1, lambda=grid)
lasso.coef = predict (out, type="coefficients", s=bestlam )
lasso.coef
```

(Intercept)	AtBat	Hits	HmRun
1.27479059	-0.05497143	2.18034583	0.00000000
RBI	Walks	Years	CAtBat
0.00000000	2.29192406	-0.33806109	0.00000000
CHmRun	CRuns	CRBI	CWalks
0.02825013	0.21628385	0.41712537	0.00000000
DivisionW	PutOuts	Assists	Errors
-116.16755870	0.23752385	0.00000000	-0.85629148

Example: Lasso ..

Estimates

```
lasso.coef[lasso.coef !=0]
```

(Intercept)	AtBat	Hits	Walks	
1.27479059	-0.05497143	2.18034583	2.29192406	-
CHmRun	CRuns	CRBI	LeagueN	
0.02825013	0.21628385	0.41712537	20.28615023	-11
PutOuts	Errors			
0.23752385	-0.85629148			

```
length(lasso.coef) # number of coefficients
```

```
[1] 20
```

```
length(lasso.coef[lasso.coef !=0]) # number of non-zero coefficients
```

```
[1] 12
```